

# DSFBA: Data Structures and Subsetting

*Data Science for Business Analytics*

Thibault Vatter

Department of Statistics, Columbia University

09/16/2020

- 1 Data structures
- 2 Atomic vectors
- 3 Attributes
- 4 S3 atomic vectors
- 5 Lists
- 6 Data frames and tibbles
- 7 Subsetting

- Type the following into your console:

```
# Create a vector in R
x <- c(5, 29, 13, 87)
x
#> [1]  5 29 13 87
```

- Two important ideas:
  - ▶ Commenting (we will come back to this)
  - ▶ Assignment
    - The `<-` symbol means assign `x` the value `c(5, 29, 13, 87)`.
    - Could use `=` instead of `<-` but this is discouraged.
    - All assignments take the same form: `object_name <- value`.
    - `c()` means "concatenate".
    - Type `x` into the console to print its assignment.

- Type the following into your console:

```
# Create a vector in R
x <- c(5, 29, 13, 87)
x
#> [1] 5 29 13 87
```

- Note: the [1] tells us that 5 is the first element of the vector.

```
# Create a vector in R
x <- 1:50
x
#> [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21
#> [22] 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42
#> [43] 43 44 45 46 47 48 49 50
```

	Homogeneous	Heterogeneous
1d	Atomic vector	List
2d	Matrix	Data frame
nd	Array	

- Almost all other objects are built upon these foundations.
- R has no 0-dimensional, or scalar types.
- Best way to understand what data structures any object is composed of is `str()` (short for structure).

```
x <- c(5, 29, 13, 87)
str(x)
#> num [1:4] 5 29 13 87
```

- Two flavors:
  - ▶ atomic vectors,
  - ▶ lists.
- Three common properties:
  - ▶ Type, `typeof()`, what it is.
  - ▶ Length, `length()`, how many elements it contains.
  - ▶ Attributes, `attributes()`, additional arbitrary metadata.
- Main difference: elements of an atomic vector must be the same type, whereas those of a list can have different types.

1 Data structures

2 Atomic vectors


3 Attributes

4 S3 atomic vectors

5 Lists

6 Data frames and tibbles

7 Subsetting



- Four primary types of atomic vectors: logical, integer, double, and character (which contains strings).
- Integer and double vectors are known as numeric vectors.
- There are two rare types: complex and raw (won't be discussed further).



Special syntax to create an individual value, AKA a **scalar**:

- **Logicals:**

- ▶ In full (TRUE or FALSE),
- ▶ Abbreviated (T or F).

- **Doubles:**

- ▶ Decimal (0.1234), scientific (1.23e4), or hexadecimal (0xcafe) form.
- ▶ Special values unique to doubles: Inf, -Inf, and NaN (not a number).

- **Integers:**

- ▶ Similar to doubles but
  - must be followed by L (1234L, 1e4L, or 0xcafeL),
  - and can not contain fractional values.


- **Strings:**

- ▶ Surrounded by " ("hi") or ' ('bye').
- ▶ Special characters escaped with \; see ?Quotes for details.

To create longer vectors from shorter ones, use `c()`:

```
lgl_var <- c(TRUE, FALSE)
int_var <- c(1L, 6L, 10L)
dbl_var <- c(1, 2.5, 4.5)
chr_var <- c("these are", "some strings")
```

Depicting vectors as connected rectangles:



- With atomic vectors, `c()` returns atomic vectors (i.e., flattens):

```
c(c(1, 2), c(3, 4))  
#> [1] 1 2 3 4
```

- Determine the type and length of a vector with `typeof()` and `length()`:

```
typeof(lgl_var)  
#> [1] "logical"  
typeof(int_var)  
#> [1] "integer"  
typeof(dbl_var)  
#> [1] "double"  
typeof(chr_var)  
#> [1] "character"
```

- Represented with NA (short for not applicable/available).
- Missing values tend to be infectious:

```
NA > 5
#> [1] NA
10 * NA
#> [1] NA
!NA
#> [1] NA
```

- Exception: when some identity holds for all possible inputs. . .

```
NA ~ 0
#> [1] 1
NA | TRUE
#> [1] TRUE
NA & FALSE
#> [1] FALSE
```

- Propagation of missingness leads to a common mistake:

```
x <- c(NA, 5, NA, 10)
x == NA
#> [1] NA NA NA NA
```

- Instead, use `is.na()`:

```
is.na(x)
#> [1] TRUE FALSE TRUE FALSE
```

- Closely related to vectors.
- Special because it has a unique type, is always length zero, and can't have any attributes.

```
typeof(NULL)
#> [1] "NULL"

length(NULL)
#> [1] 0

x <- NULL
attr(x, "y") <- 1
#> Error in attr(x, "y") <- 1: attempt to set an attribute on NULL
```

- Can test for NULLs with `is.null()`:

```
is.null(NULL)
#> [1] TRUE
```

- NULL commonly represents
  - ▶ an absent vector.
    - For example, NULL is often used as a default function argument.
    - Contrast this with NA, which indicates that an *element* of a vector is absent.
  - ▶ an empty vector (a vector of length zero) of arbitrary type.

```
c()  
#> NULL  
c(NULL, NULL)  
#> NULL  
c(NULL, 1:3)  
#> [1] 1 2 3
```

- If you're familiar with SQL, you'll know about relational NULL, but the database NULL is actually equivalent to R's NA.

- Test if a vector is of a given type with `is.*()`, but be careful:
  - ▶ `is.logical()`, `is.integer()`, `is.double()`, and `is.character()` do what you might expect.
  - ▶ Avoid `is.vector()`, `is.atomic()`, and `is.numeric()` or carefully read the documentation.
- For atomic vectors:
  - ▶ Type is a property of the entire vector (all elements of the same type).
  - ▶ When combining different types: **coercion** in a fixed order (character → double → integer → logical).

```
str(c("a", 1))  
#> chr [1:2] "a" "1"
```



- Often happens automatically:
  - ▶ Most mathematical functions (+, log, etc.) coerce to numeric.
  - ▶ Useful for logical vectors because TRUE/FALSE become 1/0.

```
x <- c(FALSE, FALSE, TRUE)
as.numeric(x)
#> [1] 0 0 1
c(sum(x), mean(x)) # Total number of TRUEs and proportion that are TRUE
#> [1] 1.000 0.333
```

- Additionally:
  - ▶ Deliberately coerce by using as.\*() (as.logical(), as.integer(), as.double(), or as.character()).
  - ▶ Failed coercion of strings → warning and missing value.

```
as.integer(c("1", "1.5", "a"))
#> Warning: NAs introduced by coercion
#> [1] 1 1 NA
```

1 Data structures

2 Atomic vectors

**3 Attributes**

4 S3 atomic vectors

5 Lists

6 Data frames and tibbles

7 Subsetting

How about matrices, arrays, factors, or date-times?


- Built on top of atomic vectors by adding attributes.
- In the next few slides:
  - ▶ The `dim` attribute to make matrices and arrays.
  - ▶ The `class` attribute to create “S3” vectors, including factors, dates, and date-times.

- Similar to name-value pairs attaching metadata to an object.
- Attributes can be retrieved/modified
  - ▶ individually with `attr()`,
  - ▶ or “En masse” with `attributes()/structure()`.

```
a <- 1:3
attr(a, "x") <- "abcdef"
attr(a, "x")
#> [1] "abcdef"

attr(a, "y") <- 4:6
str(attributes(a))
#> List of 2
#> $ x: chr "abcdef"
#> $ y: int [1:3] 4 5 6

# Or equivalently
a <- structure(
  1:3,
  x = "abcdef",
  y = 4:6
)
```



- Attributes should generally be thought of as ephemeral.
- For example, most attributes are lost by most operations:

```
attributes(a[1])  
#> NULL  
attributes(sum(a))  
#> NULL
```

- There are only two attributes that are routinely preserved:
  - ▶ `names`, a character vector giving each element a name.
  - ▶ `dim`, short for dimensions, an integer vector, used to turn vectors into matrices or arrays.
- To preserve other attributes, need to create your own S3 class!

- You can name a vector in three ways:

```
# When creating it  
x <- c(a = 1, b = 2, c = 3)  
  
# By assigning a character vector to names()  
x <- 1:3  
names(x) <- c("a", "b", "c")  
  
# Inline, with setNames()  
x <- setNames(1:3, c("a", "b", "c"))
```

- Avoid `attr(x, "names")` (more typing and less readable).
- Remove names with `unname(x)` or `names(x) <- NULL`.

- The `dim` attribute allow a vector allows it to behave like a 2-dimensional **matrix** or a multi-dimensional **array**.
- Most important feature: multidimensional subsetting, which we'll see later.
- Create matrices and arrays with `matrix()`:

```
# Two scalar arguments specify row and column sizes
a <- matrix(1:6, nrow = 2, ncol = 3)
a
#>      [,1] [,2] [,3]
#> [1,]    1    3    5
#> [2,]    2    4    6
```

## ■ Or arrays with array():

```
# One vector argument to describe all dimensions
b <- array(1:12, c(2, 3, 2))
b
#> , , 1
#>
#>      [,1] [,2] [,3]
#> [1,]    1    3    5
#> [2,]    2    4    6
#>
#> , , 2
#>
#>      [,1] [,2] [,3]
#> [1,]    7    9   11
#> [2,]    8   10   12
```



- Alternatively, use the assignment form of `dim()`:

```
# You can also modify an object in place by setting dim()
```

```
c <- 1:6
```

```
dim(c) <- c(3, 2)
```

```
c
```

```
#>      [,1] [,2]
```

```
#> [1,]    1    4
```

```
#> [2,]    2    5
```

```
#> [3,]    3    6
```

- Functions for working with vectors, matrices and arrays:

Vector	Matrix	Array
<code>names()</code>	<code>rownames()</code> , <code>colnames()</code>	<code>dimnames()</code>
<code>length()</code>	<code>nrow()</code> , <code>ncol()</code>	<code>dim()</code>
<code>c()</code>	<code>rbind()</code> , <code>cbind()</code>	<code>abind::abind()</code>
—	<code>t()</code>	<code>aperm()</code>
<code>is.null(dim(x))</code>	<code>is.matrix()</code>	<code>is.array()</code>

- A vector without a `dim` attribute set is often thought of as 1-dimensional, but actually has `NULL` dimensions.
- You also can have matrices with a single row or single column, or arrays with a single dimension:
  - ▶ They may print similarly, but will behave differently.
  - ▶ The differences aren't too important, but it's useful to know they exist in case you get strange output from a function.
  - ▶ As always, use `str()` to reveal the differences.

```
str(1:3)                                # 1d vector
#> int [1:3] 1 2 3
str(matrix(1:3, ncol = 1)) # column vector
#> int [1:3, 1] 1 2 3
str(matrix(1:3, nrow = 1)) # row vector
#> int [1, 1:3] 1 2 3
str(array(1:3, 3))              # "array" vector
#> int [1:3(1d)] 1 2 3
```

1 Data structures

2 Atomic vectors

3 Attributes


**4 S3 atomic vectors**

5 Lists

6 Data frames and tibbles

7 Subsetting


- One of the most important attributes is `class`.
  - ▶ Turns an object into an **S3 object** (which behaves differently when passed to a **generic** function).
  - ▶ Every S3 object
    - is built on top of a base type,
    - stores additional information in other attributes.
- In the next few slides, three important classes in R:
  - ▶ Categorical data (values come from a fixed set of levels): **factor** vectors.
  - ▶ Dates (day resolution): **Date** vectors.
  - ▶ Date-times (second or sub-second resolution): **POSIXct** vectors.



- A vector that can contain only predefined values.
- Used to store categorical data.
- Built on top of an integer vector with two attributes:
  - ▶ a class (defines a behavior different from integer vectors),
  - ▶ and levels (defines the set of allowed values).

```
x <- factor(c("a", "b", "b", "a"))
x
#> [1] a b b a
#> Levels: a b

typeof(x)
#> [1] "integer"
attributes(x)
#> $levels
#> [1] "a" "b"
#>
#> $class
#> [1] "factor"
```



- Useful when you know the set of possible values but they're not all present in a given dataset.
- When tabulating a factor you'll get counts of all categories, even unobserved ones:

```
sex_char <- c("m", "m", "m")
table(sex_char)
#> sex_char
#> m
#> 3

sex_factor <- factor(sex_char, levels = c("m", "f"))
table(sex_factor)
#> sex_factor
#> m f
#> 3 0
```

## ■ Ordered factors:

- ▶ Behave like regular factors, but the order of the levels is meaningful (e.g., low, medium, high)
- ▶ This property is automatically leveraged by some modelling/visualisation functions.

```
grade <- ordered(c("b", "b", "a", "c"), levels = c("c", "b", "a"))  
grade  
#> [1] b b a c  
#> Levels: c < b < a
```

## ■ While factors look like character vectors, be careful:

- ▶ Some string methods (like `gsub()` and `grep1()`) will automatically coerce factors to strings.
- ▶ Others (like `nchar()`) will throw an error.
- ▶ Still others will (like `c()`) use the underlying integer values.
- ▶ Best to explicitly convert factors to character vectors if you need string-like behavior.

- In base R:
  - ▶ Before R 4.0, factors used to be everywhere because many functions (e.g. `read.csv()/data.frame()`) automatically converted character vectors to factors.
  - ▶ Suboptimal because there's no way to know the set of all possible levels or their correct order!
- The tidyverse:
  - ▶ Never automatically coerces characters to factors.
  - ▶ Provides the `forcats` package specifically for working with factors.
  - ▶ More on that later.



- Built on top of double vectors.
- A class Date and no other attributes.

```
today <- Sys.Date()
```

```
typeof(today)
#> [1] "double"
attributes(today)
#> $class
#> [1] "Date"
```

- Value of the double = the number of days since 1970-01-01<sup>1</sup>:

```
date <- as.Date("1970-02-01")
unclass(date)
#> [1] 31
```

---

<sup>1</sup>Known as the Unix Epoch.

- Two ways of storing this information: POSIXct, and POSIXlt.
- Odd names:
  - ▶ “POSIX” is short for “Portable Operating System Interface”,
  - ▶ “ct” stands for calendar time (time\_t in C),
  - ▶ and “lt” for local time (struct tm type in C).
- Focus on POSIXct (the simplest):
  - ▶ Built on top of a double vector.
  - ▶ Value = number of seconds since 1970-01-01.

```
now_ct <- as.POSIXct("2018-08-01 22:00", tz = "UTC")
now_ct
#> [1] "2018-08-01 22:00:00 UTC"

typeof(now_ct)
#> [1] "double"
attributes(now_ct)
#> $class
#> [1] "POSIXct" "POSIXt"
#>
#> $tzone
#> [1] "UTC"
```

- The `tzzone` attribute:
  - ▶ Controls only how the formatting; not the represented instant.
  - ▶ The time is not printed if it is midnight.

```
structure(now_ct, tzzone = "Asia/Tokyo")  
#> [1] "2018-08-02 07:00:00 JST"  
structure(now_ct, tzzone = "America/New_York")  
#> [1] "2018-08-01 18:00:00 EDT"  
structure(now_ct, tzzone = "Australia/Lord_Howe")  
#> [1] "2018-08-02 08:30:00 +1030"  
structure(now_ct, tzzone = "Europe/Paris")  
#> [1] "2018-08-02 CEST"
```

1 Data structures

2 Atomic vectors

3 Attributes

4 S3 atomic vectors

**5 Lists**

6 Data frames and tibbles

7 Subsetting


- A step up in complexity from atomic vectors.
- Each element can be any type.
- Construct lists with `list()`.

```
l1 <- list(  
  1:3,  
  "a",  
  c(TRUE, FALSE, TRUE),  
  c(2.3, 5.9)  
)  
  
typeof(l1)  
#> [1] "list"  
str(l1)  
#> List of 4  
#> $ : int [1:3] 1 2 3  
#> $ : chr "a"  
#> $ : logi [1:3] TRUE FALSE TRUE  
#> $ : num [1:2] 2.3 5.9
```

1	2	3	"a"	TRUE	FALSE	TRUE	2.3	5.9
---	---	---	-----	------	-------	------	-----	-----


- Sometimes called **recursive** vectors:

```
13 <- list(list(list(1)))  
str(13)  
#> List of 1  
#> $ :List of 1  
#> ..$ :List of 1  
#> ...$ : num 1
```



- `c()` will combine several lists into one:

```
14 <- list(list(1, 2), c(3, 4))  
15 <- c(list(1, 2), c(3, 4))  
str(14)  
#> List of 2  
#> $ :List of 2  
#> ..$ : num 1  
#> ..$ : num 2  
#> $ : num [1:2] 3 4  
str(15)  
#> List of 4  
#> $ : num 1  
#> $ : num 2  
#> $ : num 3  
#> $ : num 4
```



- The `typeof()` a list is `list`.
- Test for a list with `is.list()`.
- Coerce to a list with `as.list()`.

```
list(1:3)
#> [[1]]
#> [1] 1 2 3
as.list(1:3)
#> [[1]]
#> [1] 1
#>
#> [[2]]
#> [1] 2
#>
#> [[3]]
#> [1] 3
```

- Flatten a list into an atomic vector with `unlist()`, but rules are...
  - ▶ complex & not well documented :(


- 1 Data structures
- 2 Atomic vectors
- 3 Attributes
- 4 S3 atomic vectors
- 5 Lists
- 6 Data frames and tibbles**
- 7 Subsetting



- The two most important S3 vectors built on top of lists.
- If you do data analysis in R, you'll use them.
- A data frame is a named list of vectors with attributes for (column) names, row.names, and its class, data.frame.

```
df1 <- data.frame(x = 1:3, y = letters[1:3])
typeof(df1)
#> [1] "list"

attributes(df1)
#> $names
#> [1] "x" "y"
#>
#> $class
#> [1] "data.frame"
#>
#> $row.names
#> [1] 1 2 3
```



- Similar to a list, but with an additional constraint:
  - ▶ The length of each of its vectors must be the same.
  - ▶ “Rectangular structure”:
    - Share properties of both matrices and lists.
    - Has `rownames()/colnames()` & its `names()` are the column names.
    - Has `nrow()/ncol()` & its `length()` is the number of columns.
- Data frames:
  - ▶ One of the biggest and most important ideas in R!
  - ▶ ..... but
    - 20 years have passed since their creation.
    - Lead to the creation of the tibble, a modern version.

- Provided by the `tibble` package.
- **Main difference: lazy (do less) & surly (complain more).**
- Technically:
  - ▶ Share the same structure as `data.frame`.
  - ▶ Only difference is that the class vector includes `tbl_df`.
  - ▶ Allows tibbles to behave differently.

```
library(tibble)

df2 <- tibble(x = 1:3, y = letters[1:3])
typeof(df2)
#> [1] "list"

attributes(df2)
#> $names
#> [1] "x" "y"
#>
#> $row.names
#> [1] 1 2 3
#>
#> $class
#> [1] "tbl_df"      "tbl"        "data.frame"
```

- Supply name-vector pairs to `data.frame()` or `tibble()`.

```
df <- data.frame(  
  x = 1:3,  
  y = c("a", "b", "c")  
)
```

```
df2 <- tibble(  
  x = 1:3,  
  y = c("a", "b", "c")  
)
```

```
str(df)  
#> 'data.frame':   3 obs. of  2 variables:  
#> $ x: int  1 2 3  
#> $ y: chr  "a" "b" "c"  
  
str(df2)  
#> tibble [3 x 2] (S3: tbl_df/tbl/data.frame)  
#> $ x: int [1:3] 1 2 3  
#> $ y: chr [1:3] "a" "b" "c"
```

- Next few slides: some of the differences between the two.
  - ▶ Non-syntactic names.
  - ▶ Recycling shorter inputs.
  - ▶ Variables created during construction.
  - ▶ Printing.

- Strict rules about what constitutes a valid name.
  - ▶ **Syntactic** names consist of letters<sup>2</sup>, digits, . and \_ but can't begin with \_ or a digit.
  - ▶ Additionally, can't use any of the **reserved words** like TRUE, NULL, if, and function (see the complete list in ?Reserved).
- A name that doesn't follow these rules is **non-syntactic**.

```
_abc <- 1  
#> Error: unexpected input in "_"
```

```
if <- 10  
#> Error: unexpected assignment in "if <="
```

---

<sup>2</sup>what constitutes a letter is determined by your current locale, avoid this by sticking to ASCII characters (i.e. A-Z) as much as possible.

- To override these rules and use any name:

```
`_abc` <- 1  
`_abc`  
#> [1] 1  
  
`if` <- 10  
`if`  
#> [1] 10
```

- Don't deliberately create but understand such names:
  - ▶ You'll come across them, e.g. with data created outside of R.
- In data frames and tibbles:

```
names(data.frame(`1` = 1))  
#> [1] "X1"  
  
names(data.frame(`1` = 1, check.names = FALSE))  
#> [1] "1"  
  
names(tibble(`1` = 1))  
#> [1] "1"
```

- Both `data.frame()` and `tibble()` recycle shorter inputs, but
  - ▶ data frames automatically recycle columns that are an integer multiple of the longest column,
  - ▶ tibbles will only recycle vectors of length one.

```
data.frame(x = 1:4, y = 1:2)
#>   x y
#> 1 1 1
#> 2 2 2
#> 3 3 1
#> 4 4 2
data.frame(x = 1:4, y = 1:3)
#> Error in data.frame(x = 1:4, y = 1:3): arguments imply differing
#> number of rows: 4, 3
```

- Both `data.frame()` and `tibble()` recycle shorter inputs, but
  - ▶ data frames automatically recycle columns that are an integer multiple of the longest column,
  - ▶ tibbles will only recycle vectors of length one.

```
tibble(x = 1:4, y = 1)
```

```
#> # A tibble: 4 x 2
```

```
#>       x     y
```

```
#>   <int> <dbl>
```

```
#> 1     1     1
```

```
#> 2     2     1
```

```
#> 3     3     1
```

```
#> 4     4     1
```

```
tibble(x = 1:4, y = 1:2)
```

```
#> Error: Tibble columns must have compatible sizes.
```

```
#> * Size 4: Existing data.
```

```
#> * Size 2: Column `y`.
```

```
#> i Only values of size one are recycled.
```



- `tibble()` allows you to refer to variables created during construction:

```
tibble(  
  x = 1:3,  
  y = x * 2  
)  
#> # A tibble: 3 x 2  
#>       x     y  
#>   <int> <dbl>  
#> 1     1     2  
#> 2     2     4  
#> 3     3     6
```

(Inputs are evaluated left-to-right.)

iris

#>	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
#> 1	5.1	3.5	1.4	0.2	setosa
#> 2	4.9	3.0	1.4	0.2	setosa
#> 3	4.7	3.2	1.3	0.2	setosa
#> 4	4.6	3.1	1.5	0.2	setosa
#> 5	5.0	3.6	1.4	0.2	setosa
#> 6	5.4	3.9	1.7	0.4	setosa
#> 7	4.6	3.4	1.4	0.3	setosa
#> 8	5.0	3.4	1.5	0.2	setosa
#> 9	4.4	2.9	1.4	0.2	setosa
#> 10	4.9	3.1	1.5	0.1	setosa
#> 11	5.4	3.7	1.5	0.2	setosa
#> 12	4.8	3.4	1.6	0.2	setosa
#> 13	4.8	3.0	1.4	0.1	setosa
#> 14	4.3	3.0	1.1	0.1	setosa
#> 15	5.8	4.0	1.2	0.2	setosa
#> 16	5.7	4.4	1.5	0.4	setosa
#> 17	5.4	3.9	1.3	0.4	setosa
#> 18	5.1	3.5	1.4	0.3	setosa
#> 19	5.7	3.8	1.7	0.3	setosa
#> 20	5.1	3.8	1.5	0.3	setosa
#> 21	5.4	3.4	1.7	0.2	setosa

```
dplyr::starwars
#> # A tibble: 87 x 14
#>   name height mass hair_color skin_color eye_color birth_year
#>   <chr>   <int> <dbl> <chr>      <chr>      <chr>      <dbl>
#> 1 Luke~   172    77 blond     fair       blue       19
#> 2 C-3P0   167    75 <NA>     gold      yellow    112
#> 3 R2-D2    96    32 <NA>     white, bl~ red       33
#> 4 Dart~   202   136 none      white     yellow    41.9
#> 5 Leia~   150    49 brown     light     brown     19
#> 6 Owen~   178   120 brown, gr~ light     blue     52
#> 7 Beru~   165    75 brown     light     blue     47
#> 8 R5-D4    97    32 <NA>     white, red red      NA
#> 9 Bigg~   183    84 black     light     brown     24
#> 10 Obi~   182    77 auburn, w~ fair      blue-gray 57
#> # ... with 77 more rows, and 7 more variables: sex <chr>,
#> #   gender <chr>, homeworld <chr>, species <chr>, films <list>,
#> #   vehicles <list>, starships <list>
```

- Only the first 10 rows + the columns that fit on screen.
- Each column is labelled with its abbreviated type.
- Wide columns are truncated.
- In RStudio, color highlights important information.

- To check if an object is a data frame or tibble:

```
is.data.frame(df1)
#> [1] TRUE
is.data.frame(df2)
#> [1] TRUE
```

- Typically, it should not matter if you have a tibble or data frame, but if you need to be certain:

```
is_tibble(df1)
#> [1] FALSE
is_tibble(df2)
#> [1] TRUE
```

- Coerce an object to a data frame or tibble with `as.data.frame()` or `as_tibble()`.

- 1 Data structures
- 2 Atomic vectors
- 3 Attributes
- 4 S3 atomic vectors
- 5 Lists
- 6 Data frames and tibbles
- 7 Subsetting**

- R's subsetting operators are fast and powerful.
  - ▶ Allows to succinctly perform complex operations in a way that few other languages can match.
  - ▶ Easy to learn but hard to master because of a number of interrelated concepts:
    - Six ways to subset atomic vectors.
    - Three subsetting operators, `[`, `[[`, and `$`.
    - The operators interact differently with different vector types.
    - Subsetting can be combined with assignment.
- Subsetting is a natural complement to `str()`:
  - ▶ `str()` shows the pieces of any object (its structure).
  - ▶ Subsetting pulls out the pieces that you're interested in.
- Outline:
  - ▶ Selecting multiple elements with `[`.
  - ▶ Selecting a single element with `[[` and `$`.
  - ▶ Subsetting and assignment.

- We'll look at the following vector:

```
x <- c(2.1, 4.2, 3.3, 5.4)
```

- Note that the number after the decimal point represents the original position in the vector.
- There are six things that you can use to subset a vector:
  - ▶ Positive integers.
  - ▶ Negative integers.
  - ▶ Logical vectors.
  - ▶ Nothing.
  - ▶ Zero.
  - ▶ Character vectors.

## ■ Positive integers return elements at the specified positions:

```
x[c(3, 1)]  
#> [1] 3.3 2.1  
x[order(x)]  
#> [1] 2.1 3.3 4.2 5.4  
  
x[c(1, 1)] # Duplicate indices will duplicate values  
#> [1] 2.1 2.1  
  
x[c(2.1, 2.9)] # Real numbers are silently truncated to integers  
#> [1] 4.2 4.2
```

## ■ Negative integers exclude elements at the specified positions:

```
x[-c(3, 1)]  
#> [1] 4.2 5.4
```

## ■ Can't mix positive and negative integers in a single subset:

```
x[c(-1, 2)]  
#> Error in x[c(-1, 2)]: only 0's may be mixed with negative subscripts
```



- **Logical vectors** select elements where the corresponding logical value is TRUE (probably the most useful):

```
x[c(TRUE, TRUE, FALSE, FALSE)]  
#> [1] 2.1 4.2  
x[x > 3]  
#> [1] 4.2 3.3 5.4
```

- In `x[y]`, what happens if `x` and `y` are different lengths?
  - ▶ Behavior controlled by the **recycling rules** with the shorter recycled to the length of the longer.
  - ▶ Convenient and easy to understand when `x` OR `y` is length one, but avoid for other lengths because of inconsistencies in base R.

```
x[c(TRUE, FALSE)]  
#> [1] 2.1 3.3  
# Equivalent to  
x[c(TRUE, FALSE, TRUE, FALSE)]  
#> [1] 2.1 3.3
```

- **Nothing** returns the original vector (not useful for 1D vectors, but important for matrices, data frames, and arrays:

```
x[]  
#> [1] 2.1 4.2 3.3 5.4
```

- **Zero** returns a zero-length vector (not usually done on purpose):

```
x[0]  
#> numeric(0)
```

- If the vector is named, you can also use **character vectors** to return elements with matching names:

```
(y <- setNames(x, letters[1:4]))
```

```
#>   a    b    c    d
```

```
#> 2.1 4.2 3.3 5.4
```

```
y[c("d", "c", "a")]
```

```
#>   d    c    a
```

```
#> 5.4 3.3 2.1
```

```
# Like integer indices, you can repeat indices
```

```
y[c("a", "a", "a")]
```

```
#>   a    a    a
```

```
#> 2.1 2.1 2.1
```

```
# When subsetting with [, names are always matched exactly
```

```
z <- c(abc = 1, def = 2)
```

```
z[c("a", "d")]
```

```
#> <NA> <NA>
```

```
#>   NA   NA
```

- Note that a missing value in the index always yields a missing value in the output:

```
x[c(TRUE, TRUE, NA, FALSE)]  
#> [1] 2.1 4.2 NA
```

- Factors are not treated specially when subsetting:
  - ▶ Subsetting will use the underlying integer vector, not the character levels.
  - ▶ Typically unexpected, so avoid!

```
y[factor("b")]  
#> a  
#> 2.1
```

- Exactly as for atomic vectors.
- Using [ always returns a list; [[ and \$ (see in a few slides), lets you pull out elements of a list.

- Subset higher-dimensional structures in three ways:
  - ▶ With multiple vectors.
  - ▶ With a single vector.
  - ▶ With a matrix.
- The most common way:
  - ▶ Supply a 1D index for each dimension, separated by a comma.
  - ▶ Blank subsetting is now useful!

```
a <- matrix(1:9, nrow = 3)
colnames(a) <- c("A", "B", "C")
a[1:2, ]
#>      A B C
#> [1,] 1 4 7
#> [2,] 2 5 8
a[c(TRUE, FALSE, TRUE), c("B", "A")]
#>      B A
#> [1,] 4 1
#> [2,] 6 3
a[0, -2]
#>      A C
```

- By default, [ simplifies the results to the lowest possible dimensionality.
  - ▶ For example, both of the following expressions return 1D vectors.
  - ▶ You'll learn how to avoid “dropping” dimensions later.

```
a[1, ]  
#> A B C  
#> 1 4 7  
a[1, 1]  
#> A  
#> 1
```

- Can subset them with a vector as if they were 1D.
- Note that arrays in R are stored in column-major order:

```
vals <- outer(1:5, 1:5, FUN = "paste", sep = ",")
```

```
vals
```

```
#>      [,1] [,2] [,3] [,4] [,5]  
#> [1,] "1,1" "1,2" "1,3" "1,4" "1,5"  
#> [2,] "2,1" "2,2" "2,3" "2,4" "2,5"  
#> [3,] "3,1" "3,2" "3,3" "3,4" "3,5"  
#> [4,] "4,1" "4,2" "4,3" "4,4" "4,5"  
#> [5,] "5,1" "5,2" "5,3" "5,4" "5,5"
```

```
vals[c(4, 15)]
```

```
#> [1] "4,1" "5,3"
```



- Can also subset higher-dimensional data structures with an integer matrix (or, if named, a character matrix).
  - ▶ Each row in the matrix specifies the location of one value.
  - ▶ Each column corresponds to a dimension in the array.
  - ▶ E.g., use a 2 column matrix to subset a matrix, a 3 column matrix to subset a 3D array, etc.
  - ▶ The result is a vector of values.

```
select <- matrix(ncol = 2, byrow = TRUE, c(
  1, 1,
  3, 1,
  2, 4
))

vals[select]
#> [1] "1,1" "3,1" "2,4"
```

- Characteristics of both lists and matrices.
- When subsetting with a single index:
  - ▶ Behave like lists and index the columns.
  - ▶ E.g. `df[1:2]` selects the first two columns.
- When subsetting with two indices:
  - ▶ Behave like matrices.
  - ▶ E.g. `df[1:3, ]` selects the first three *rows* (and all columns)<sup>3</sup>.

```
df <- data.frame(x = 1:3, y = 3:1, z = letters[1:3])
```

```
df[df$x == 2, ]
```

```
#>   x y z
```

```
#> 2 2 2 b
```

```
df[c(1, 3), ]
```

```
#>   x y z
```

```
#> 1 1 3 a
```

```
#> 3 3 1 c
```

---

<sup>3</sup>In Python `df[1:3, 1:2]` would select three columns and two rows.

- Two ways to select columns from a data frame:

```
# Like a list
df[c("x", "z")]
#>   x z
#> 1 1 a
#> 2 2 b
#> 3 3 c
# Like a matrix
df[, c("x", "z")]
#>   x z
#> 1 1 a
#> 2 2 b
#> 3 3 c
```

- Important difference if you select a single column:
  - ▶ Matrix subsetting simplifies by default.
  - ▶ List subsetting does not.

```
str(df[, "x"])  
#> int [1:3] 1 2 3  
str(df["x"])  
#> 'data.frame': 3 obs. of 1 variable:  
#> $ x: int 1 2 3
```

- Subsetting a tibble with [ always returns a tibble:

```
df <- tibble::tibble(x = 1:3, y = 3:1, z = letters[1:3])  
  
str(df["x"])  
#> tibble [3 x 1] (S3: tbl_df/tbl/data.frame)  
#> $ x: int [1:3] 1 2 3  
str(df[, "x"])  
#> tibble [3 x 1] (S3: tbl_df/tbl/data.frame)  
#> $ x: int [1:3] 1 2 3
```

- For matrices and arrays, dimensions with length 1 are dropped:

```
a <- matrix(1:4, nrow = 2)
str(a[1, ])
#> int [1:2] 1 3

str(a[1, , drop = FALSE])
#> int [1, 1:2] 1 3
```

- Data frames with a single column returns just that column:

```
df <- data.frame(a = 1:2, b = 1:2)
str(df[, "a"])
#> int [1:2] 1 2

str(df[, "a", drop = FALSE])
#> 'data.frame': 2 obs. of 1 variable:
#> $ a: int 1 2
```

- The default `drop = TRUE` is a common source of bugs:
  - ▶ Your code with a dataset with multiple columns works.
  - ▶ Six months later, you use it with a single column dataset and it fails with a mystifying error.
  - ▶ **Always use 'drop = FALSE' when subsetting a 2D object!**
  - ▶ Tibbles default to `drop = FALSE` and `[` always returns a tibble.
- Factor subsetting also has a `drop` argument:
  - ▶ Controls whether or not levels (rather than dimensions) are preserved defaults to `FALSE`.
  - ▶ When using `drop = TRUE`, use a character vector instead.

```
z <- factor(c("a", "b"))
z[1]
#> [1] a
#> Levels: a b
z[1, drop = TRUE]
#> [1] a
#> Levels: a
```

The other two subsetting operators:

- `[[` is used for extracting single items.
- `x$y` is a useful shorthand for `x[["y"]]`.


- `[[` is most important when working with lists because subsetting a list with `[` always returns a smaller list.

*If list  $x$  is a train carrying objects, then  $x[[5]]$  is the object in car 5;  $x[4:6]$  is a train of cars 4-6.*

— @RLangTip, <https://twitter.com/RLangTip/status/268375867468681216>


- Use this metaphor to make a simple list:

```
x <- list(1:3, "a", 4:6)
```






- When extracting a single element, you have two options:
  - ▶ Create a smaller train, i.e., fewer carriages, with `[.`
  - ▶ Extract the contents of a particular carriage with `[[.`



- When extracting multiple (or even zero!) elements, you have to make a smaller train.



- Shorthand operator:
  - ▶ `x$y` is roughly equivalent to `x[["y"]]`.
  - ▶ Often used to access variables in a data frame.
  - ▶ E.g., `mtcars$cyl` or `diamonds$carat`.
- One common mistake with `$`:

```
var <- "cyl"
# Doesn't work - mtcars$var translated to mtcars[["var"]]
mtcars$var
#> NULL

# Instead use [[
mtcars[[var]]
#> [1] 6 6 4 6 8 6 8 4 4 6 6 8 8 8 8 8 8 4 4 4 4 8 8 8 8 4 4 4 8 6 8 4
```

- The one important difference between \$ and [[ is (left-to-right) partial matching:

```
x <- list(abc = 1)
x$a
#> [1] 1
x[["a"]]
#> NULL
```

- To avoid this, the following is highly recommended:

```
options(warnPartialMatchDollar = TRUE)
x$a
#> Warning in x$a: partial match of 'a' to 'abc'
#> [1] 1
```

(For data frames, you can also avoid this problem by using tibbles, which never do partial matching.)

- Data frames have two undesirable subsetting behaviors.
  - ▶ When you subset columns with `df[, vars]`:
    - Returns a vector if `vars` selects one variable.
    - Otherwise, returns a data frame.
    - Frequent unless you use `drop = FALSE`.
  - ▶ When extracting a single column with `df$x`:
    - If there is no column `x`, selects any variable that starts with `x`.
    - If no variable starts with `x`, returns `NULL`.
    - Easy to select the wrong variable/a variable that doesn't exist.
- Tibbles tweak these behaviors:
  - ▶ `[` always returns a tibble.
  - ▶ `$` doesn't do partial matching and warns if it can't find a variable (makes tibbles surly).

```
df1 <- data.frame(xyz = "a")
str(df1$x)
#> chr "a"
```

```
df2 <- tibble(xyz = "a")
str(df2$x)
#> Warning: Unknown or uninitialised column: `x`.
#> NULL
```

- Subsetting operators can be combined with assignment.
  - ▶ Modifies selected values of an input vector
  - ▶ Called subassignment.
- The basic form is `x[i] <- value`:

```
x <- 1:5  
x[c(1, 2)] <- c(101, 102)  
x  
#> [1] 101 102 3 4 5
```

- Recommendation:
  - ▶ Make sure that `length(value)` is the same as `length(x[i])`,
  - ▶ and that `i` is unique.
  - ▶ Otherwise, you'll end-up in recycling hell.

## ■ Subsetting lists with NULL

- ▶ `x[[i]] <- NULL` removes a component.
- ▶ To add a literal NULL, use `x[i] <- list(NULL)`.

```
x <- list(a = 1, b = 2)
x[["b"]] <- NULL
str(x)
#> List of 1
#> $ a: num 1
```

```
y <- list(a = 1, b = 2)
y[["b"]] <- list(NULL)
str(y)
#> List of 2
#> $ a: num 1
#> $ b: NULL
```

## ■ Subsetting with nothing can be useful with assignment

- ▶ Preserves the structure of the original object.
- ▶ Compare the following two expressions.

```
mtcars[] <- lapply(mtcars, as.integer)
is.data.frame(mtcars)
#> [1] TRUE
```

```
mtcars <- lapply(mtcars, as.integer)
is.data.frame(mtcars)
#> [1] FALSE
```